ConcordAI: A Multi-Agent Conversational System for Conflict Mediation

Gunda Rama Praneetha
dept. of Computing Technologies
SRM Institute of Science and Technology
Kattankulathur, Tamil Nadu, India
rg0337@srmist.edu.in

Archit Gupta
dept. of Computing Technologies
SRM Institute of Science and Technology
Kattankulathur, Tamil Nadu, India
ag0595@srmist.edu.in

Rounaq Moin
dept. of Computing Technologies
SRM Institute of Science and Technology
Kattankulathur, Tamil Nadu, India
rm3062@srmist.edu.in

Dr. A. Maria Nancy
dept. of Computing Technologies
SRM Institute of Science and Technology
Kattankulathur, Tamil Nadu, India
marianaa@srmist.edu.in

Saswat Singh
ZS, Management Consulting
Services Company
Bengaluru, India
saswat.singh@zs.com

***Abstract*—In this work, we propose ConcordAI, a multi-agent conversational pipeline for mediating between individuals who hold opposing viewpoints. The system is designed to generate balanced and empathetic responses that help users interpret each other's perspectives more constructively. The architecture consists of a perspective extraction agent for determining the relationship between the users' stances, a sentiment and intent agent for estimating emotional signals using fine-grained emotion scores, a reconciliation response agent for generating mediation-oriented responses using the identified stances and emotional context, and a safety moderation module that filters unsafe responses using a predefined toxicity threshold. Automated evaluation was conducted on a dataset of 90 conflict scenarios created for the proposed system. The evaluation reports the accuracy of stance and sentiment prediction for both users. Additionally, the system was evaluated through a blind human evaluation with 56 participants, where each participant was presented with responses generated by ConcordAI and GPT-5.2 for the same conflict scenarios and asked to select the response they preferred. The results show that 82.1% of participants preferred responses generated by the proposed system, thus indicating that it reliably identifies perspective and emotional direction and can serve as a practical building block for real-world conflict mediation and assistance tools.**

***Index Terms*—Multi-Agent Systems, Large Language Models (LLMs), Stance Detection, Intent Extraction, Conflict Resolution, Empathetic Response Generation, Conversational AI.**

I. Introduction

The increasing use of AI-based conversational systems has significantly altered the nature of user interactions with computational platforms. However, such platforms often face challenges in addressing situations characterized by divergent user viewpoints. For example, in online platforms such as discussion forums, customer support platforms, and digital working spaces, users often articulate a wide range of viewpoints regarding a wide variety of issues, ranging from social issues to daily divergences.

Conventional conversational AI platforms face challenges in addressing situations involving multiple users with different opinions. For example, traditional AI systems often generate neutral responses that fail to address emotions expressed by users, resulting in users feeling ignored and misinterpreted. In some cases, these systems incline towards a single viewpoint due to bias or design limitation, which displays unfairness. The primary challenge associated with these conversational AI platforms is that they are generally designed for one-to-one interactions, which often fail to address situations with multiple users.

To overcome these limitations, this work presents a multi-agent conversational system called ConcordAI. It is intended to produce responses in a mediated form in a scenario with opposing viewpoints. Rather than neutral or biased response, it processes the input based on the user's stance and emotional context before producing a response. It follows a modular pipeline in which various agents process different components of the dialogue and perform conflict mediation.

The proposed system objectives are primarily to make sure the responses represent both sides of the dialogue and not favor one side of the argument. Secondly, the responses should address the emotional context of the dialogue so that the participants feel their concerns have been heard. Also, the proposed system should remain clear and logical in its responses while promoting a constructive dialogue between the participants, and ensure that the responses are checked using a safety assessment tool to avoid the generation of any harmful responses.

The system consists of four cooperating agents, arranged in a pipeline, as illustrated in Fig. 1. (1) Perspective Extractor Agent: Takes the initial user inputs and identifies their respective stances and key claims. (2) Sentiment and Intent Agent: Given the original text, this agent identifies emotions and thus deduces underlying intentions. (3) Reconciliation Response Agent: Combines structured stance and intent information to generate a candidate response that is balanced and empathetic. (4) Safety Moderation Agent: Evaluates for toxicity, harmful

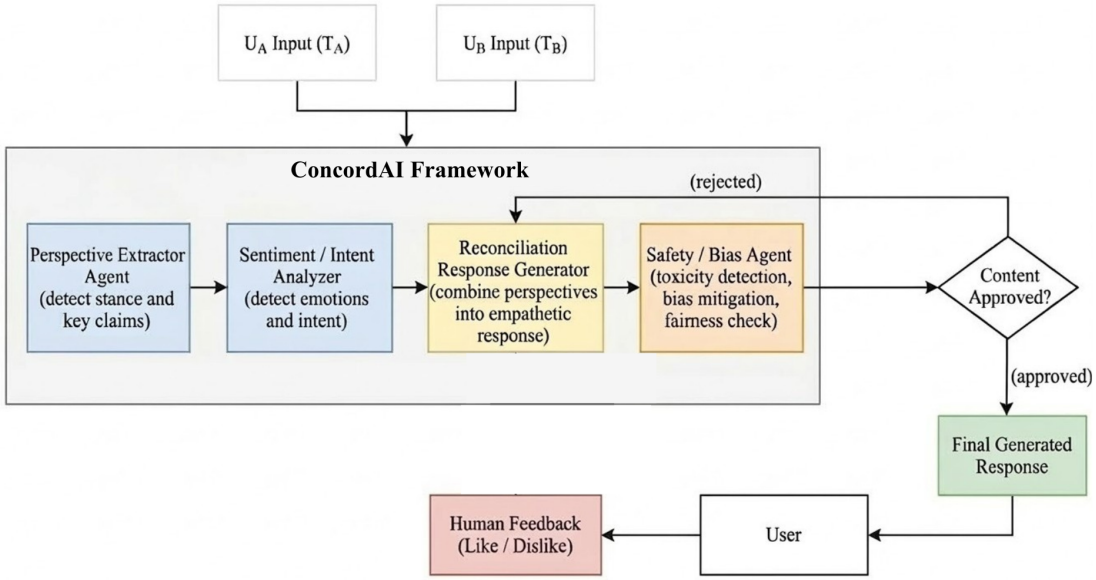

Fig. 1. Conceptual Diagram of ConcordAI Pipeline

content, and fairness; after this, the approved responses are returned to users, while rejected responses set off iterative refinement by the reconciliation agent.

II. Related Work

The creation of AI systems designed to mediate and resolve user conflicts draws upon several related research areas. These include empathetic conversational AI, multi-agent systems for coordination and conflict resolution, and methods for detecting and reducing bias, as well as perspective analysis. Empathetic responses are essential for effective mediation. Recent research has explored how to integrate commonsense reasoning and intent detection into dialogue systems [1], along with models that use self–other awareness to balance empathy with neutrality [2]. Additionally, surveys of empathetic agents have studied current architectures and highlighted ongoing challenges, such as empathy transfer across different domains [3]. Similarly, prompt design strategies have been studied to improve emotional awareness and the quality of responses in these systems [4]. These developments provide the basis for the stance - intent detection and response-generation components used in our system.

Multi-agent architecture offers a framework for managing viewpoints. A study on the effectiveness of multi-agent debate systems indicates that a structured flow of information between agents can enhance the results of the reasoning process [5]. AI systems for resolving conflicts have been shown to be effective in enhancing satisfaction levels among the stakeholders in a collaborative environment [6]. AI4PCR systems illustrate the application of AI for training users in constructive communication and mediation techniques [7]. In addition, logical approaches based on structured argumentation have shown results for reconciling disagreements in a balanced way [8]. Motivated by these research areas, our system proposes a multi-agent pipeline that processes conflicting inputs through sequential reasoning stages.

Ensuring bias reduction and fairness are important considerations to conversational AI systems [9]. Ethical and cultural considerations with respect to the design of agent-based systems are formal strategies that are advanced [10]. Besides that, the regulatory and governance strategies emphasize the importance of fairness evaluation in the development of an AI system [4]. Several benchmarks have been proposed to assess the bias of language models. These benchmarks are referred to as CHBias [11] and BiasAsker [12]. At the same time, the development of the ProsocialDialog dataset [13] has promoted research and development of dialogue systems that emphasize safety and social awareness.

Accurate representation of viewpoints is fundamental. Perspectivist natural language processing (NLP) approaches emphasize the explicit modeling of varied perspectives [14]. Stance detection techniques, such as SocialPET [15], demonstrate feasibility even in environments with limited resources. Frameworks such as ECON [16] utilize consistency modeling to manage contradictory evidence. Our Perspective Extraction Agent is constructed upon these established methods. Furthermore, domain-specific applications provide additional validation for AI's utility in conflict resolution, including its application in workplace grievance procedures [17] and the use of emotion recognition to enhance dispute resolution outcomes [18]. Our approach is also informed by broader frameworks designed for resolving value conflicts within the context of public AI governance [19].

III. System Architecture

The architecture of the ConcordAI system is structured in a way that it provides balanced and empathetic responses to the conflicting viewpoints of two users. Instead of providing

TABLE I
AGENTS IN THE CONCORDAI PIPELINE

Agent	Role	Model / Tool	Training Data	Input	Output
Perspective Extraction	Detects stance relationships between user statements	DeBERTa-v3 NLI cross-encoder	SNLI, MultiNLI	User texts (T_A, T_B)	`perspective_json`
Sentiment and Intent	Extracts emotion scores, intent, and valence	RoBERTa (samlowe/roberta-base-go-emotions)	GoEmotions dataset	User texts (T_A, T_B)	`sentiment_json`
Reconciliation Response	Generates mediation-oriented responses	Mistral-7B-Instruct (via Ollama)	Mistral pretraining data	User texts (T_A, T_B) + perspective_json + sentiment_json	$R_{candidate}$
Safety Moderation	Evaluates generated responses for toxicity and safety	Detoxify	Jigsaw Unintended Bias	$R_{candidate}$	Safety decision JSON

common and neutral responses, the system aims to analyze the stance and emotional content of the users' statements before providing mediation-oriented responses. The ConcordAI system uses a modular multi-agent pipeline architecture. Here, different agents perform different tasks and exchange structured outputs. This architecture helps the mediation process at different stages to focus on certain aspects of the interaction while maintaining the flow of information.

The architecture of the system is based on various design principles. Perspective preservation ensures that the viewpoints of both users continue to be represented throughout the processing pipeline so that the final response $R_{\text{candidate}}$ takes the viewpoints expressed in the conversation into account. Sequential specialization structures the architecture as a modular multi-agent pipeline, where each agent carries out a specific task of analysis or generation and produces structured outputs that act as input to the next processing steps. Safety-aware processing adds a moderation step to the pipeline that checks the responses produced and filters out any harmful or toxic responses. Finally, the goal of empathetic response synthesis is to produce responses that address the viewpoints of both users and promote constructive dialogue.

The system has four cooperating agents (as stated in TABLE I) that are arranged in sequence in a pipeline, as shown in Fig.2. These agents are Perspective Extraction Agent, Sentiment and Intent Agent, Reconciliation Response Agent, and Safety Moderation Agent. These components receive data from the previous system stage and then produce outputs in JSON format.

i. *Perspective Extraction Agent's* goal is to determine the relationship between both user stances. The model used in this stage is cross-encoder/nli-deberta-v3-large, which has been trained on NLI datasets such as SNLI and MultiNLI. This transformer-based Natural Language Inference (NLI) is to analyze how the statements of the two users relate to each other. The model takes textual inputs T_A and T_B from the users, processes them and assignes either of these three NLI labels: contradiction, neutral, or entailment. They represent the structured interpretation of how the users' statements interact. The NLI labels are then mapped to stance representations to describe if the two users support or oppose each other. This provides a structured JSON format $perspective_json$ of the scenario, which is used in the later stages of the pipeline.

Input: Text data from User A (T_A) and User B (T_B).

Output: A structured JSON $perspective_json$ object containing the predicted stance relationships for both users. The output follows the JSON structure shown in Listing 1.

```
{
  "user_A_perspectives": [
    {
      "actor": "user_A",
      "stance": "<support | oppose>",
      "key_claims": ["<user_A_statement>"]
    }
  ],
  "user_B_perspectives": [
    {
      "actor": "user_B",
      "stance": "<support | oppose>",
      "key_claims": ["<user_B_statement>"]
    }
  ]
}
```

Listing 1. Perspective Extraction Agent Output Format

ii. *Sentiment and Intent Agent's* goal is to grasp the emotional intent of both user inputs T_A and T_B. Emotional intent is crucial to be able to reply to emotions in a constructive mediation response. The agent employs a transformer-based model called samlowe/roberta-base-go-emotions, which has been pretrained on Google's GoEmotions dataset. The probabilities are calculated over 28 unique emotion classes, hence facilitating the system's ability to recognize emotional cues from user input. An emotion is considered relevant if its probability is greater than 0.3. Additionally, the emotion with the greatest probability is considered to be the user's intent which is stored in a JSON format $sentiment_json$ including the valence and emotion scores . A simple valence is computed as 1 to grasp the emotional intent of user input.

$$valence = joy - anger \tag{1}$$

This provides a fundamental emotional state of whether the message reflects a more positive or negative emotional state.

Input: Text data from User A (T_A) and User B (T_B).

Output: object containing dominant intent, and valence scores for both users. The $sentiment_json$ follows the structure shown in Listing 2.

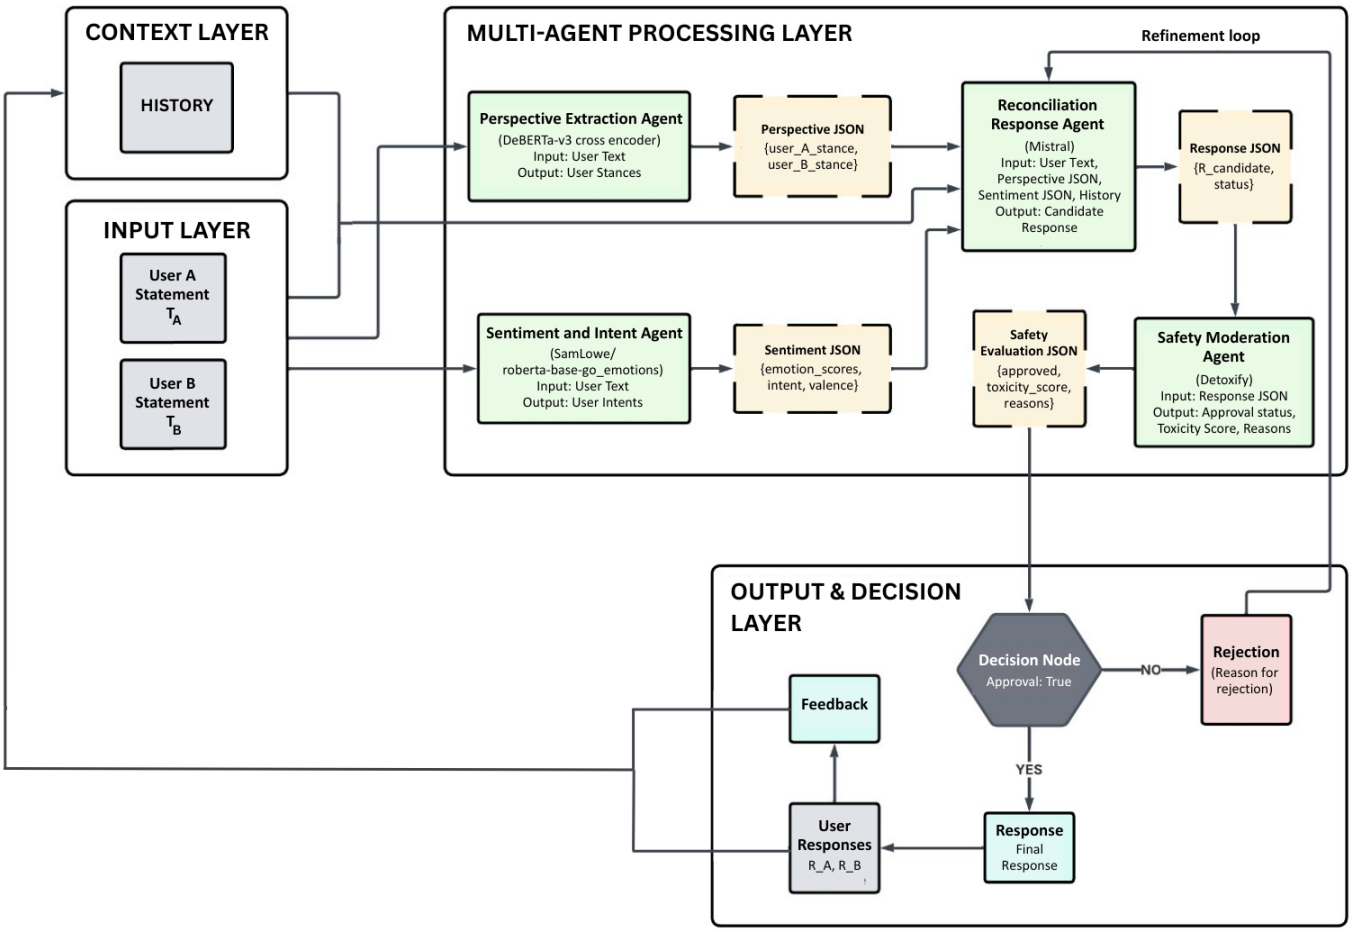

Fig. 2. Architecture Diagram of ConcordAI

```
{
  "user_A_sentiment": {
    "valence": <float>,
    "emotion_scores": {
      "<emotion_label>": <float>,
      "<emotion_label>": <float>
    },
    "intent": "<dominant_emotion>"
  },

  "user_B_sentiment": {
    "valence": <float>,
    "emotion_scores": {
      "<emotion_label>": <float>,
      "<emotion_label>": <float>
    },
    "intent": "<dominant_emotion>"
  }
}
```

Listing 2. Sentiment and Intent Agent Output Format

iii. *Reconciliation Response Agent's* goal is to generate an empathetic response $R_{\text{candidate}}$ for each user based on stance and intent detected from earlier stages. The agent uses the Mistral-7B-Instruct language model, executed locally through Ollama, to generate responses. A structured prompt is constructed using both user statements T_A and T_B, the stance relationship between them, the emotional intent behind the statement and response instructions for mediation. It commands the model to generate a response acknowledging both viewpoints and maintain a respectful and constructive tone. The generated output is further parsed and separated into segments with two responses R_A, R_B and $status$ of the conversation with labels: Escalated, Stable, Improved, Resolved. When a conversation is continued, the history containing user texts and responses is also provided as context, and based on this, the agent generates $R_{candidate}$ and $status$ of the conversation.

Input: Text data from User A (T_A) and User B (T_B), $perspective_json$, $sentiment_json$ and History H.

Output: A JSON object with the response text $R_{candidate}$.

```
{
  "candidate_response": "<full_mediation_message>",
  "status": "<escalating | stable | improving | resolved>"
}
```

Listing 3. Reconciliation Response Agent Output Format

iv. *Safety Moderation Agent* is responsible for a thorough check of the generated response $R_{candidate}$ on issues related to toxicity, bias, fairness, and consistency. The system uses the Detoxify toxicity classifier, which is trained on the Jigsaw Unintended Bias dataset, to estimate the toxicity level of the generated response. If the toxicity score exceeds a threshold

of 0.3, the response is unsafe. If the response satisfies this constraint, the response is $approved$ and sent to users. If it does not pass the constraint, the response is rejected and the previous agent is prompted for refinement of the response.

Input: Generated response $R_{candidate}$ in JSON format.

Output: A JSON object with the field indicating whether the response is approved or rejected, and reasons for rejection (if any) based on the toxicity evaluation. The output follows the JSON structure shown in Listing 4.

```
{
  "approved": true | false,
  "toxicity_score": <float>,
  "reason": "<rejection_reason_if_any>"
}
```

Listing 4. Safety Moderation Agent Output Format

Algorithm 1 Multi-Agent Conflict Mediation Pipeline

Require: User A text T_A, User B text T_B, conversation history H

Ensure: Mediation responses R_A, R_B and conversation status

Generate $conversation_id$
Receive user inputs T_A and T_B
$perspective_json \leftarrow PerspectiveAgent(T_A, T_B)$
$sentiment_json \leftarrow SentimentAgent(T_A, T_B)$
$R_{candidate} \leftarrow ReconciliationAgent(T_A, T_B,$
$\quad perspective_json, sentiment_json,$
$\quad history = None)$
$approved \leftarrow SafetyAgent(R_{candidate}).approved$
while $approved = false$ **do**
  $R_{candidate} \leftarrow ReconciliationAgent(T_A, T_B,$
  $\quad perspective_json, sentiment_json,$
  $\quad history = H)$
  $approved \leftarrow SafetyAgent(R_{candidate}).approved$
end while
Extract responses R_A, R_B and $status$ from $R_{candidate}$
return R_A, R_B, $status$

The agents operate in a centralized control pipeline, (as shown in Algorithm 1) where the coordination of information flow between agents is carried out. At the start of the interaction, a unique $conversation_id$ is generated, enabling the tracking of the interaction. Subsequently, the textual data provided by both users is fed into the Perspective Extraction Agent, where the relationship between the two statements is established. The statements are also given to the Sentiment and Intent Agent for analyzing intent. The output generated by both agents $perspective_json$ and $sentiment_json$ is sent to the Reconciliation Response Agent, where a candidate response $R_candidate$ is generated. This response is checked by the Safety Moderation Agent for the presence of harmful content. If the generated response is $approved$, it is sent to the users. If it fails, a new response is generated.

The communication between agents is carried out in a standardized JSON format, allowing for clear interfaces between agents. This enables the extension and modification of individual agents without affecting the system.

IV. Experimental Setup

The proposed system has been implemented in Python, and it is based on a modular multi-agent system with a central controller. The controller takes two user statements and optional conversation history, and it generates mediation responses via the proposed agents, as described in Section III. The proposed agents communicate with each other in JSON format, allowing modular communication between these agents during execution.

To test our proposed system in conflict mediation scenarios, we have developed our own dataset, containing 90 two-party conversation scenarios, acting as ground truth. Each scenario consists of two user statements that reflect a short conflict or tension in conversation. The dataset we're proposing includes a variety of scenarios. These encompass toxic comments, instances of misunderstanding, situations involving power imbalances, expressions of emotional frustration, boundary-setting behaviors, and cultural missteps.

Each scenario is provided with ground truth information, where the relationship between the participants in the conversation is annotated. In other words, the ground truth information includes stance information, where the stance is either support or oppose, and sentiment intent, where the sentiment intent is based on the dominant emotional tone in the statement. The ground truth information is in a structured format, denoted by JSON. A simple representation of the ground truth schema is depicted in Listing 5.

```
{
  "case_id": "<unique_scenario_identifier>",
  "category": "<conflict_type>",

  "user_a_text": "<statement_from_user_A>",
  "user_b_text": "<statement_from_user_B>",

  "user_a_perspective": {
    "stance": "<support | oppose>"
  },

  "user_b_perspective": {
    "stance": "<support | oppose>"
  },

  "user_a_sentiment": {
    "sentiment": "<positive | negative>"
  },

  "user_b_sentiment": {
    "sentiment": "<positive | negative>"
  }
}
```

Listing 5. Ground Truth Schema

For the automated evaluation, each scenario is evaluated by the system just once with an empty conversation history. The stance and sentiment intent predicted by the analytical agents are compared with the ground truth information. The metrics accuracy and recall are used to measure the performance, and these are calculated for stance and sentiment in both user statements . All experiments were conducted on a single workstation running 64-bit Windows and using the Python 3.12 environment in a virtual environment. The inference was conducted on CPU without any GPU acceleration to mimic a

lightweight deployment scenario. The same environment was used for all experiments.

V. Results and Discussion

A. Automation Evaluation

The system demonstrates strong alignment with human annotations on the curated dataset, and the results are summarized in I. Stance prediction is the most reliable component of the analytical pipeline. The Perspective Extraction agent achieves an accuracy of 0.91 for User A and 0.80 for User B, with recall values of 0.95 and 0.89, respectively. These results indicate that the model consistently distinguishes supportive and opposing conversational moves in short conflict exchanges.

The Sentiment and Intent agent achieved an accuracy of 0.80 for User A and 0.82 for User B. The corresponding recall values of 0.89 and 0.90 show that the system is able to capture the overall emotional direction of the scenarios despite the loss of nuance compared to the full GoEmotions label space.

TABLE II
Automated evaluation results on the Ground Truth.

Task	Accuracy	Recall
Stance – User A	0.91	0.95
Stance – User B	0.80	0.89
Sentiment – User A	0.80	0.89
Sentiment – User B	0.82	0.90

Overall, automated evaluation suggests that analytical components of the pipeline can reliably extract the conversational signals required for mediation, namely stance relationships and emotional intent. These signals provide the contextual information used by the response generation component to produce balanced mediation responses.

These results support the claim that the multi-agent pipeline can reliably characterize conversational perspective and emotional intent in short single-turn conflicts while remaining computationally feasible for interactive mediation systems.

B. Human Evaluation

A study was conducted to assess the performance of the proposed method in realistic disagreement scenarios by human evaluation. The study involved 56 participants. Each participant was presented with pairs of conflicting user statements together with two AI-generated responses: one produced by the proposed ConcordAI system and another generated using GPT-5.2. The responses were presented in a blind comparison format, and participants were asked to select the response they considered more effective for mediating the disagreement.

Participants evaluated which response better acknowledged both viewpoints and encouraged constructive mediation. The aggregated preference results are illustrated in Fig. 3. As a result, it is observed that 82.1% of participants preferred the proposed system for conflict mediation for given user statements, while 17.9% preferred the responses generated by GPT-5.2.

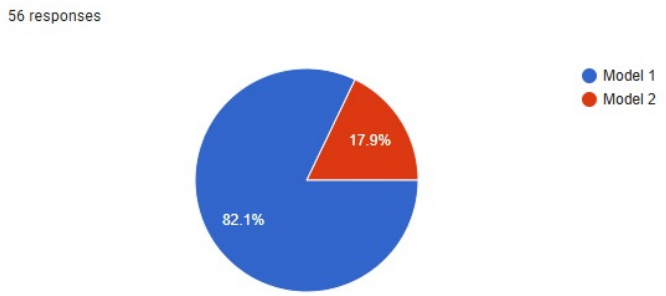

Fig. 3. Participant preference between ConcordAI and GPT-5.2 responses.

In addition to the overall preference comparison, participants also evaluated responses across several qualitative criteria, including fairness, mediation ability, clarity, and emotional awareness. The results of these additional evaluation categories are shown in Fig. 4. For fairness, the proposed system received 45 votes, compared to 11 votes for the GPT-5.2 model. For mediation ability, the proposed system received 41 votes, while the compared model received 14 votes. Similarly, the clarity in response for the proposed system received 43 votes compared to 17 votes in the other model, and 40 votes for emotional awareness were obtained compared to 15 votes for GPT-5.2.

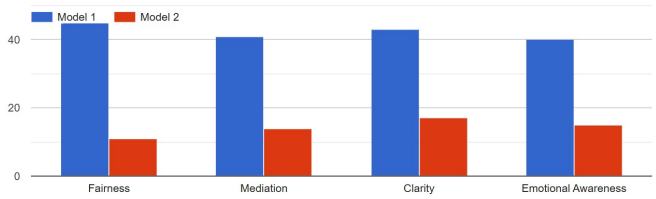

Fig. 4. Human evaluation comparison across fairness, mediation ability, clarity, and emotional awareness.

The evaluation findings suggest that the designed pipeline generates responses that participants typically find to be clearer, more balanced, and more emotionally attuned. In terms of qualitative assessment, responses generated by ConcordAI consistently received a greater number of votes compared to those produced by the compared model. Consequently, the human evaluation supports the efficacy of the proposed pipeline in managing dyadic conflicts and promoting more productive dialogue among participants.

VI. Conclusion

The present work proposes a multi-agent system named ConcordAI, which may assist in solving the problem of disagreements between different users having different views on a particular situation. It consists of four different agents: Perspective Extraction, Sentiment and Intent, Reconciliation Response, and Safety Moderation. These agents analyze the interaction and generate a response that is both empathetic and balanced. Automated experiments using the curated dataset of 90 scenarios show that the analytical agents can accurately

determine the relationships between stances and the emotional intent of the interactions.

A human evaluation of the proposed method with 56 participants was conducted to evaluate the effectiveness of the pipeline through a blind preference test with GPT-5.2. It was found that 82.1% of the participants chose responses generated by ConcordAI. These results show that the suggested workflow can effectively mediate conflicts between users with conflicting viewpoints on a particular issue.

VII. Future Work

The future work will focus on extending the adaptability of the system. One way of doing that is to include reward-based feedback mechanisms within the pipeline, which can implement reinforcement learning from human feedback (RLHF) to improve the reconciliation agent based on user preferences about fairness, clarity, and empathy. Another addition to the proposed work can be to allow for mediation for multi-party conflicting scenarios. Conflicts in the real world may involve more than two parties. Therefore, the proposed system will need to be extended to allow for more complex conversations. Finally, the proposed work can be evaluated on multi-turn conversations, since the proposed work is only evaluated on single-turn scenarios. The proposed work can also be improved by allowing the system to explore domain adaptation and culturally aware mediation mechanisms, as this can help agents capture context while maintaining the modularity of the pipeline.

References

[1] X. Wang, T. Liu, Z. Liu, and Z. Fang, “EmpCI: Empathetic response generation with common sense and empathetic intent,” Knowledge-Based Systems, vol. 283, p. 111128, 2023.

[2] W. Zhao, Y. Zhao, X. Lu, and B. Qin, “Don’t Lose Yourself! Empathetic response generation via explicit self-other awareness,” in Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL), Toronto, Canada, 2023, pp. 14602–14616.

[3] A. S. Raamkumar and Y. Yang, “Empathetic Conversational Systems: A Review of Current Advances, Gaps, and Opportunities,” IEEE Trans. Affective Computing, vol. 14, no. 4, pp. 2722-2739, 2023.

[4] Y. Zhang, C. Radishian, S. Brunswicker, D. Whitenack, and D. W. Linna Jr., “Empathetic Language in LLMs under Prompt Engineering: A Comparative Study in the Legal Field,” Procedia Computer Science, vol. 244, pp. 308-317, 2024, doi:10.1016/j.procs.2024.10.204.

[5] A. Tillmann, “Literature Review of Multi-Agent Debate for Problem-Solving,” arXiv preprint arXiv:2506.00066, May 2025.

[6] R. Aydoğan, T. Baarslag, and E. Gerding, ”Artificial Intelligence Techniques for Conflict Resolution,” Group Decision and Negotiation, vol. 30, no. 4, pp. 879–883, Aug. 2021, doi: 10.1007/s10726-021-09738-x.

[7] A. Hsu and D. Chaudhary, ”AI4PCR: Artificial intelligence for practicing conflict resolution,” Computers in Human Behavior: Artificial Humans, vol. 1, no. 1, p. 100002, 2023, doi: 10.1016/j.chbah.2023.100002.

[8] S. L. Vasileiou, A. Kumar, W. Yeoh, T. C. Son, and F. Toni, ”Dialectical Reconciliation via Structured Argumentative Dialogues,” in Proc. 21st Int. Conf. on Principles of Knowledge Representation and Reasoning (KR), 2024.

[9] S. Westerstrand, “Fairness in AI systems development: EU AI Act compliance and beyond,” Information and Software Technology, vol. 187, 2025, Art. 107864, doi:10.1016/j.infsof.2025.107864.

[10] S. G. Yaman, P. Ribeiro, A. Cavalcanti, R. Calinescu, C. Paterson, and B. Townsend, “Specification, validation and verification of social, legal, ethical, empathetic and cultural requirements for autonomous agents,” Journal of Systems and Software, vol. 214, p. 112233, 2024.

[11] J. Zhao et al., ”CHBias: Bias Evaluation and Mitigation of Chinese Conversational Language Models,” in Proc. 61st Annu. Meeting of the Assoc. for Computational Linguistics (ACL), Toronto, Canada, 2023, pp. 13539–13552, doi: 10.18653/v1/2023.acl-long.757.

[12] Y. Wan, W. Wang, P. He, J. Gu, H. Bai, and M. Lyu, ”BiasAsker: Measuring the Bias in Conversational AI System,” in Proc. 31st ACM Joint Eur. Software Eng. Conf. and Symp. on the Foundations of Software Eng. (ESEC/FSE), San Francisco, CA, USA, 2023, pp. 515–527, doi: 10.1145/3611643.3616310.

[13] H. Kim et al., ”PROSOCIALDIALOG: A Prosocial Backbone for Conversational Agents,” in Proc. 2022 Conf. on Empirical Methods in Natural Language Processing (EMNLP), Abu Dhabi, United Arab Emirates, 2022, pp. 3844–3858, doi: 10.18653/v1/2022.emnlp-main.267.

[14] S. Frenda et al., ”Perspectivist approaches to natural language processing: a survey,” Lang. Resour. Evaluation, vol. 59, no. 2, pp. 1719–1746, 2025, doi: 10.1007/s10579-024-09766-4.

[15] P. Jamadi Khiabani and A. Zubiaga, ”SocialPET: Socially Informed Pattern Exploiting Training for Few-Shot Stance Detection in Social Media,” in Proc. 2024 Joint Int. Conf. on Comput. Linguistics, Lang. Resour. and Eval. (LREC-COLING), Torino, Italy, 2024, pp. 1630–1643.

[16] J. Cheng et al., ”ECON: On the Detection and Resolution of Evidence Conflicts,” in Proc. 2024 Conf. on Empirical Methods in Natural Language Processing (EMNLP), Miami, FL, USA, 2024, pp. 7816–7831.

[17] Y. S. D. ATTHOS, K. M. Iftikhar, F. Imtiaz, F. Azim, and S. A. Bangash, ”The Role of AI in Conflict Resolution and Employee Grievance Handling in HRM,” Dialogue Social Science Review (DSSR), vol. 3, no. 3, pp. 1164–1178, 2025.

[18] N. Farooq, Z. Rafique, A. Merchant, and K. R. Narejo, ”Investigating how Artificial Intelligence AI Systems Powered by Emotional Recognition Technologies can assist in Conflict Resolution by Understanding Human Emotions and Behavior,” Review of Applied Management and Social Sciences, vol. 8, no. 1, pp. 193–205, 2025, doi: 10.47067/ramss.v8i1.449.

[19] K. de Fine Licht, “Resolving value conflicts in public AI governance: A procedural justice framework,” Government Information Quarterly, vol. 42, no. 2, Jun. 2025, Art. 102033, doi:10.1016/j.giq.2025.102033.